



Shape Optimisation with PDE Constraints

Mike Fuller (fuller@maths.ox.ac.uk)

Supervised by Dr Alberto Paganini (paganini@maths.ox.ac.uk)

12th July 2018



Abstract: Shape optimisation is typically concerned with finding the shape which is optimal in the sense that it minimises a certain cost functional while satisfying given constraints. Often we find that the functional being solved depends on the solution of a PDE constraint.

In this report we briefly discuss unconstrained optimisation methods* before focussing on a shape optimisation test case that, despite its relative simplicity, exhibits an unexpected and interesting behaviour: solving this problem via Newton's method with a *truncated* second shape derivative performs better than using full second order information. We perform numerical (and some theoretical) investigations to try to shed light on this unexpected behaviour.

*One method we will explore is the *Steepest Descent* method; an example using `gradientDescentDemo` from the Machine Learning Toolbox is pictured below the title. MATLAB code by *Roger Jang* and *MIR Lab*, available at mirlab.org/jang/matlab/toolbox/machineLearning

Contents

1	Unconstrained Optimisation	3
1.1	Setting the Scene	3
1.2	Descent Methods	3
1.3	Line Search	4
1.4	Steepest Descent	4
1.5	Newton's Method	6
2	Shape Differentiation	8
2.1	Defining the Shape Derivative	8
2.2	Shape Derivative Identities	8
3	Shape Optimisation Test Case	10
4	Conclusion	11

1 Unconstrained Optimisation

We first recap on some iterative methods for solving unconstrained optimisation problems.

1.1 Setting the Scene

Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be a function. Suppose we want to find a value x^* such that f is minimal, i.e

$$x^* \in \arg \min_x f(x).$$

We can find such a value by means of an iterative algorithm. That is, given an initial guess $x^{(0)}$ for x^* , we construct a *minimising sequence* $(x^{(k)})$ of points such that

$$f(x^{(k)}) \longrightarrow f(x^*) \quad \text{as } k \rightarrow \infty.$$

Often we only require an approximate solution, so we can reduce computational effort by choosing a tolerance $\varepsilon > 0$ such that when $f(x^{(k)}) - f(x^*) \leq \varepsilon$, we terminate the algorithm.

1.2 Descent Methods

The minimising sequences we are going to discuss in this section are all of the form

$$x^{(k+1)} = x^{(k)} + t^{(k)} \Delta x^{(k)}$$

where

- $\Delta x^{(k)}$ represents a *search direction* - which direction we travel in to approach x^* ,
- $t^{(k)} \geq 0$ represents a *step length* - how far we travel along the vector $\Delta x^{(k)}$.

Recall that our aim is to minimise f , so we would like the sequence to have the property

$$f(x^{(k+1)}) < f(x^{(k)}) \tag{1}$$

for all k , unless $x^{(k)}$ is optimal. If f is differentiable, we can Taylor expand up to first-order and write

$$f(x^{(k+1)}) = f(x^{(k)} + t^{(k)} \Delta x^{(k)}) \approx f(x^{(k)}) + t^{(k)} \mathrm{d}f(x^{(k)}, \Delta x^{(k)})$$

where $\mathrm{d}f$ denotes the *directional derivative*

$$\mathrm{d}f(x, v) := \lim_{t \rightarrow 0} \frac{f(x + tv) - f(x)}{t}.$$

Hence, if $x^{(k)}$ is not optimal (so that $t^{(k)} \neq 0$), we can ensure (1) holds up to first-order if

$$\mathrm{d}f(x^{(k)}, \Delta x^{(k)}) < 0. \tag{2}$$

A search direction $\Delta x^{(k)}$ such that (2) holds is called a *descent direction*.

A generic descent method then goes as follows:

```
1 Given starting point  $x$ 
2 repeat
3   | (1) Determine descent direction,  $\Delta x$ 
4   | (2) Choose step size,  $t > 0$ 
5   | (3) Set  $x \leftarrow x + t\Delta x$ 
6 until stopping criterion satisfied
```

At the moment this is rather vague. We have a few questions to answer:

- How do we choose the step size?
- What is the descent direction?
- When do we stop the algorithm?

1.3 Line Search

To answer the first question, we discuss two methods of deciding the step size, called *exact line search* (ELS) and *backtracking line search* (BLS). Here we assume that Δx is a descent direction.

In ELS, we choose t such that f is minimised along the ray $\{x + t\Delta x : t \geq 0\}$, i.e

$$t = \arg \min_{s \geq 0} f(x + s\Delta x).$$

It may be difficult to find t exactly, so it is generally cost effective to choose a t value such that f is only approximately minimised along the ray. For this reason we may choose to apply BLS, which generates the step size by the following algorithm:

Algorithm 1: Backtracking Line Search

```

1 Given descent direction  $\Delta x$ ,  $\alpha \in (0, 0.5)$ ,  $\beta \in (0, 1)$ 
2  $t := 1$ 
3 while  $f(x + t\Delta x) > f(x) + \alpha t \, df(x, \Delta x)$  do
4   | Set  $t \leftarrow \beta t$ 
5 end
```

The resulting value of t is chosen as the step size. The line search is called "backtracking" as it starts with unit step size, then reduces it by the factor β until the stopping criterion

$$f(x + t\Delta x) > f(x) + \alpha t \, df(x, \Delta x)$$

holds. Since Δx is a descent direction, we have $df(x, \Delta x) < 0$, so

$$f(x + t\Delta x) \approx f(x) + t \, df(x, \Delta x) < f(x) + \alpha t \, df(x, \Delta x)$$

for small enough t , demonstrating that BLS eventually terminates. The choice of parameters α, β has a noticeable but not dramatic effect on convergence (see [1]); ELS can sometimes improve convergence rate but not dramatically. It is generally not worth the computational effort of using ELS.

1.4 Steepest Descent

We can find a descent direction by analysing the Taylor expansion of $f(x + \Delta x)$ about x up to first-order. The approximation is

$$f(x + \Delta x) \approx f(x) + df(x, \Delta x)$$

so we find the direction of *steepest descent* when we minimise the directional derivative. We define a *normalised steepest descent direction* Δx_{nsd} as a descent direction of unit length that minimises $df(x, \Delta x)$, i.e

$$\Delta x_{\text{nsd}} \in \arg \min_{\|\Delta x\|=1} df(x, \Delta x)$$

for some norm $\|\cdot\|$ on \mathbb{R}^n .

To piece together the *Steepest Descent* method, it will be useful to introduce the concept of inner products and gradients, particularly when it comes to defining stopping criteria. Let V be a real vector space. An *inner product* $\langle \cdot, \cdot \rangle : V^2 \rightarrow \mathbb{R}$ is a function that satisfies the following properties:

- *Linearity*: $\langle \alpha u + \beta v, w \rangle = \alpha \langle u, w \rangle + \beta \langle v, w \rangle \quad \forall u, v, w \in V, \alpha, \beta \in \mathbb{R}$
- *Symmetry*: $\langle u, v \rangle = \langle v, u \rangle \quad \forall u, v \in V$
- *Positive definiteness*: $\langle u, u \rangle \geq 0 \quad \forall u \in V, \quad \langle u, u \rangle = 0 \iff u = 0.$

It follows from the first two properties that $\langle \cdot, \cdot \rangle$ is *bilinear*.

We call $\nabla f(x)$ the *gradient* of f at x with respect to an inner product $\langle \cdot, \cdot \rangle$ if

$$\langle \nabla f(x), v \rangle = \mathrm{d}f(x, v) \quad \forall v.$$

We occasionally write $\nabla_{\langle \cdot, \cdot \rangle} f(x)$ to make it clear what the associated inner product is.

Lemma 1.1: If a norm $\|\cdot\|$ is induced by an inner product $\langle \cdot, \cdot \rangle$, i.e $\|v\| = \sqrt{\langle v, v \rangle}$, then

$$\Delta x_{\text{nsd}} = -\frac{\nabla f(x)}{\|\nabla f(x)\|}. \quad (3)$$

Proof: We use the method of Lagrange multipliers. Let

$$L(\Delta x, \lambda) := \mathrm{d}f(x, \Delta x) + \frac{\lambda}{2}(\|\Delta x\|^2 - 1).$$

Then by definition of Δx_{nsd} ,

$$\begin{aligned} \partial_{\Delta x} L(\Delta x_{\text{nsd}}, v) &= \mathrm{d}f(x, v) + \lambda \langle \Delta x_{\text{nsd}}, v \rangle = 0 \quad \forall v \\ \iff \langle \lambda \Delta x_{\text{nsd}}, v \rangle &= -\mathrm{d}f(x, v) \quad \forall v \\ \iff \lambda \Delta x_{\text{nsd}} &= -\nabla f(x). \end{aligned}$$

Taking norms of both sides, and using the fact that $\|\Delta x_{\text{nsd}}\| = 1$, we obtain $|\lambda| = \|\nabla f(x)\|$. The result follows, using the fact that Δx_{nsd} is a steepest descent direction. \square

Geometrically speaking, we see that Δx_{nsd} is a step in the unit ball of $\|\cdot\|$ which extends farthest in the direction of $-\nabla f(x)$. It is more convenient in practice to use an unnormalised descent direction

$$\Delta x_{\text{sd}} := \|\nabla f(x)\|_* \Delta x_{\text{nsd}}$$

where $\|\cdot\|_*$ denotes the *dual norm*

$$\|v\|_* := \sup_{\|x\| \leq 1} v^T x.$$

We find that Δx_{sd} is simply the unnormalised version of (3), that is

$$\Delta x_{\text{sd}} = -\nabla_{\langle \cdot, \cdot \rangle} f(x)$$

and we arrive at the *Steepest Descent* method, which uses Δx_{sd} as a descent direction:

Algorithm 2: Steepest Descent

```

1 Given starting point  $x$ 
2 repeat
3   (1) Set  $\Delta x \leftarrow \Delta x_{\text{sd}}$ 
4   (2) Choose  $t$  via ELS/BLS
5   (3) Set  $x \leftarrow x + t\Delta x$ 
6 until stopping criterion satisfied
```

To define a stopping criterion, recall that when x^* is optimal we have

$$df(x^*, v) = 0 \quad \forall v$$

so for a sufficiently smooth f , we expect the gradient to be small for $x \approx x^*$. For this reason, we can define a stopping criterion of the form

$$\|\nabla f(x)\| \leq \varepsilon$$

for some chosen small $\varepsilon > 0$. Most practical uses of Steepest Descent check the stopping criterion after the first step.

1.5 Newton's Method

Which inner product do we use to define Δx_{sd} ? An interesting choice, which forms the basis of *Newton's Method*, is to consider the inner product induced by the second derivative

$$d^2 f(x, u, v) := \lim_{t \rightarrow 0} \frac{df(x + tv, u) - df(x, u)}{t}.$$

If it exists, we can extend our Taylor approximation up to second-order to get

$$f(x + \Delta x) \approx f(x) + df(x, \Delta x) + \frac{1}{2} d^2 f(x, \Delta x, \Delta x) =: g(\Delta x).$$

If $d^2 f(x, \cdot, \cdot)$ is positive definite then g is *strictly convex*[†], thus has a unique minimiser Δx_{nt} , called the *Newton step*. Since the second derivative exists, $d^2 f(x, \cdot, \cdot)$ is also symmetric and linear, hence an inner product. By minimality, the Newton step must satisfy

$$\begin{aligned} dg(\Delta x_{nt}, v) &= 0 \quad \forall v \\ \iff df(x, v) + d^2(x, \Delta x_{nt}, v) &= 0 \quad \forall v \\ \iff d^2(x, \Delta x_{nt}, v) &= -df(x, v) \quad \forall v \end{aligned}$$

so Δx_{nt} is the negative gradient with respect to the inner product induced by $d^2 f$, that is,

$$\Delta x_{nt} = -\nabla_{d^2 f(x, \cdot, \cdot)} f(x).$$

If f is quadratic, then $f(\Delta x_{nt}) = g(\Delta x)$ exactly, so $x + \Delta x_{nt}$ is the *exact* minimiser of f . Hence it should also be a very good estimate for x^* if f is approximately quadratic. Since f is assumed to be twice continuously differentiable, the quadratic model of f is very accurate for x near x^* , so $x + \Delta x_{nt}$ is a very good estimate for x^* .

Before we introduce *Newton's method*, we construct a stopping criterion using the *Newton decrement*

$$\lambda(x) := \sqrt{df(x, \Delta x_{nt})}.$$

Lemma 1.2: $\lambda(x)^2/2$ estimates the error $f(x) - f(x^*)$, based on the quadratic approximation of f .

Proof: The error based on the quadratic approximation is

$$\begin{aligned} f(x) - \inf_y g(y) &= f(x) - g(\Delta x_{nt}) = f(x) - \left(f(x) + \cancel{df(x, \Delta x_{nt})} + \frac{1}{2} d^2 f(x, \Delta x_{nt}, \Delta x_{nt}) \right) \\ &= f(x) - \left(f(x) - \frac{1}{2} df(x, \Delta x_{nt}) \right) \\ &= \frac{1}{2} \lambda(x)^2. \end{aligned}$$

□

[†]A function g is said to be *strictly convex* if the line segment connecting any two distinct points on the surface of g lies strictly above g , except at the endpoints.

We use this as a stopping criterion for Newton's Method, which goes as follows:

Algorithm 3: Newton's Method

```

1 Given starting point  $x$ , tolerance  $\varepsilon > 0$ 
2 repeat
3   (1) Set  $\Delta x \leftarrow \Delta x_{\text{nt}}$ 
4   (2) Set  $\lambda \leftarrow \lambda(x)$ 
5   (3) Stop if  $\lambda^2/2 \leq \varepsilon$ 
6   (4) Choose  $t$  via BLS
7   (5) Set  $x \leftarrow x + t\Delta x$ 

```

Observe that this is more or less a general descent method, with the difference that the stopping criterion is checked after computing Δx_{nt} , rather than after updating the value of x .

Newton's Method has *quadratic convergence* near x^* (see [1]). Roughly speaking, this means that the number of correct digits doubles after each iteration. Newton's Method also scales with the size of the problem: its performance in \mathbb{R}^{10000} is similar to problems in \mathbb{R}^{10} say, with only a modest increase in the number of iterations.

A pitfall of Newton's Method is the cost of computing Δx_{nt} , which requires solving a system of linear equations. *Quasi-Newton* methods require less cost to form the search direction, sharing some advantages of Newton's Method such as rapid convergence near x^* , but we will not discuss such methods here.

2 Shape Differentiation

Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be a function. Provided f is smooth, we can differentiate to find its minimum values; in this section we extend the notion of differentiation of functions to *shape differentiation* of domains, with the aim of using these shape derivatives to find the domain Ω^* in a collection of admissible shapes \mathcal{U}_{ad} which minimises the cost functional $\mathcal{J}[\Omega] := \int_{\Omega} f \, dx$. That is,

$$\Omega^* \in \arg \min_{\Omega \in \mathcal{U}_{\text{ad}}} \mathcal{J}[\Omega].$$

2.1 Defining the Shape Derivative

Let $T : \mathbb{R}^n \rightarrow \mathbb{R}^n$ be a sufficiently smooth vector field. Then

$$\mathcal{J}[T(\Omega)] = \int_{T(\Omega)} f \, dx = \int_{\Omega} (f \circ T) |\det DT| \, dx$$

where DT denotes the *Jacobian* of T . Note that $T(\Omega) \neq \Omega$ in general, and similarly $\mathcal{J}[T(\Omega)] \neq \mathcal{J}[\Omega]$.

Furthermore let $V : \mathbb{R}^n \rightarrow \mathbb{R}^n$. We adopt the notation

$$\left\langle \frac{d}{dT} g[T], V \right\rangle := \lim_{t \rightarrow 0} \frac{g[T + tV] - g[T]}{t}$$

noting that this is different to the inner product notation used in Section 1.4. We then define the *shape derivative* of \mathcal{J} at Ω in the direction V as

$$d\mathcal{J}(\Omega, V) := \left\langle \frac{d}{dT} \mathcal{J}[T(\Omega)], V \right\rangle \Big|_{T=I}$$

where I is the identity map. Therefore,

$$\begin{aligned} d\mathcal{J}(\Omega, V) &= \left\langle \frac{d}{dT} \left(\int_{\Omega} (f \circ T) |\det DT| \, dx \right), V \right\rangle \Big|_{T=I} \\ &= \int_{\Omega} \left\langle \frac{d}{dT} ((f \circ T) \det DT), V \right\rangle \Big|_{T=I} \, dx. \end{aligned}$$

2.2 Shape Derivative Identities

We prove two identities of the shape derivative using the above.

Theorem 2.1: Let $\nabla \cdot V$ denote the *divergence* of V . Then

$$d\mathcal{J}(\Omega, V) = \int_{\Omega} \nabla f \cdot V + f \nabla \cdot V \, dx.$$

Proof: By the product rule,

$$d\mathcal{J}(\Omega, V) = \int_{\Omega} \left\{ \left\langle \frac{d}{dT} (f \circ T), V \right\rangle \det DT + \left\langle \frac{d}{dT} \det DT, V \right\rangle (f \circ T) \right\} \Big|_{T=I} \, dx.$$

From the first term, we have

$$\left\langle \frac{d}{dT} (f \circ T), V \right\rangle = \frac{d}{dt} (f \circ (T + tV)) = \nabla(f \circ T) \cdot V$$

so that

$$\left\langle \frac{d}{dT} (f \circ T), V \right\rangle \det DT \Big|_{T=I} = \nabla f \cdot V.$$

From the second term, using Jacobi's formula we have

$$\begin{aligned}
\left\langle \frac{d}{dT} \det DT, V \right\rangle &= \lim_{t \rightarrow 0} \frac{\det D(T + tV) - \det DT}{t} \\
&= \lim_{t \rightarrow 0} \frac{\det(DT + tDV) - \det DT}{t} \\
&= \lim_{t \rightarrow 0} \frac{(\det DT + t \operatorname{Tr}(\operatorname{adj}(DT)DV) + O(t^2)) - \det DT}{t} \\
&= \operatorname{Tr}(\operatorname{adj}(DT)DV)
\end{aligned}$$

where $\operatorname{Tr}(A)$, $\operatorname{adj}(A)$ denote the *trace* and *adjugate* of A respectively. Hence

$$\left\langle \frac{d}{dT} \det DT, V \right\rangle (f \circ T) \Big|_{T=I} = \operatorname{Tr}(DV)f = f \nabla \cdot V$$

and the result follows. \square

Theorem 2.2: We have the identity

$$d\mathcal{J}(\Omega, V) = \int_{\partial\Omega} f(V \cdot n) dS$$

where n is the outward pointing unit normal of the boundary $\partial\Omega$.

Proof: Using the Divergence Theorem,

$$\begin{aligned}
d\mathcal{J}(\Omega, V) &= \int_{\Omega} \nabla f \cdot V + f \nabla \cdot V dx \\
&= \int_{\Omega} \nabla \cdot (Vf) dx \\
&= \int_{\partial\Omega} (Vf) \cdot n dS \\
&= \int_{\partial\Omega} f(V \cdot n) dS
\end{aligned}$$

where the second equality follows from vector calculus. \square

3 Shape Optimisation Test Case

4 Conclusion

References

- [1] Stephen Boyd and Lieven Vandenberghe. *Convex Optimisation*. Cambridge University Press, 2004. web.stanford.edu/~boyd/cvxbook/bv_cvxbook.pdf